

# JVM and Code Generation For Java

Wuu Yang

wuuyang@cs.nctu.edu.tw

National Chiao-Tung University, Taiwan, R.O.C.

first draft: May 7, 2004

current version: December 28, 2011

This set of notes is adapted from JVM Spec, Chap 7, Compiling for  
the Java Virtual Machine.

Copyright ©December 28, 2011 by Wu Yang. All rights reserved.

# Contents

<b>1</b>	<b>Java Virtual Machines</b>	<b>4</b>
<b>2</b>	<b>Bytecode</b>	<b>21</b>
<b>3</b>	<b>Sample Java Code</b>	<b>76</b>

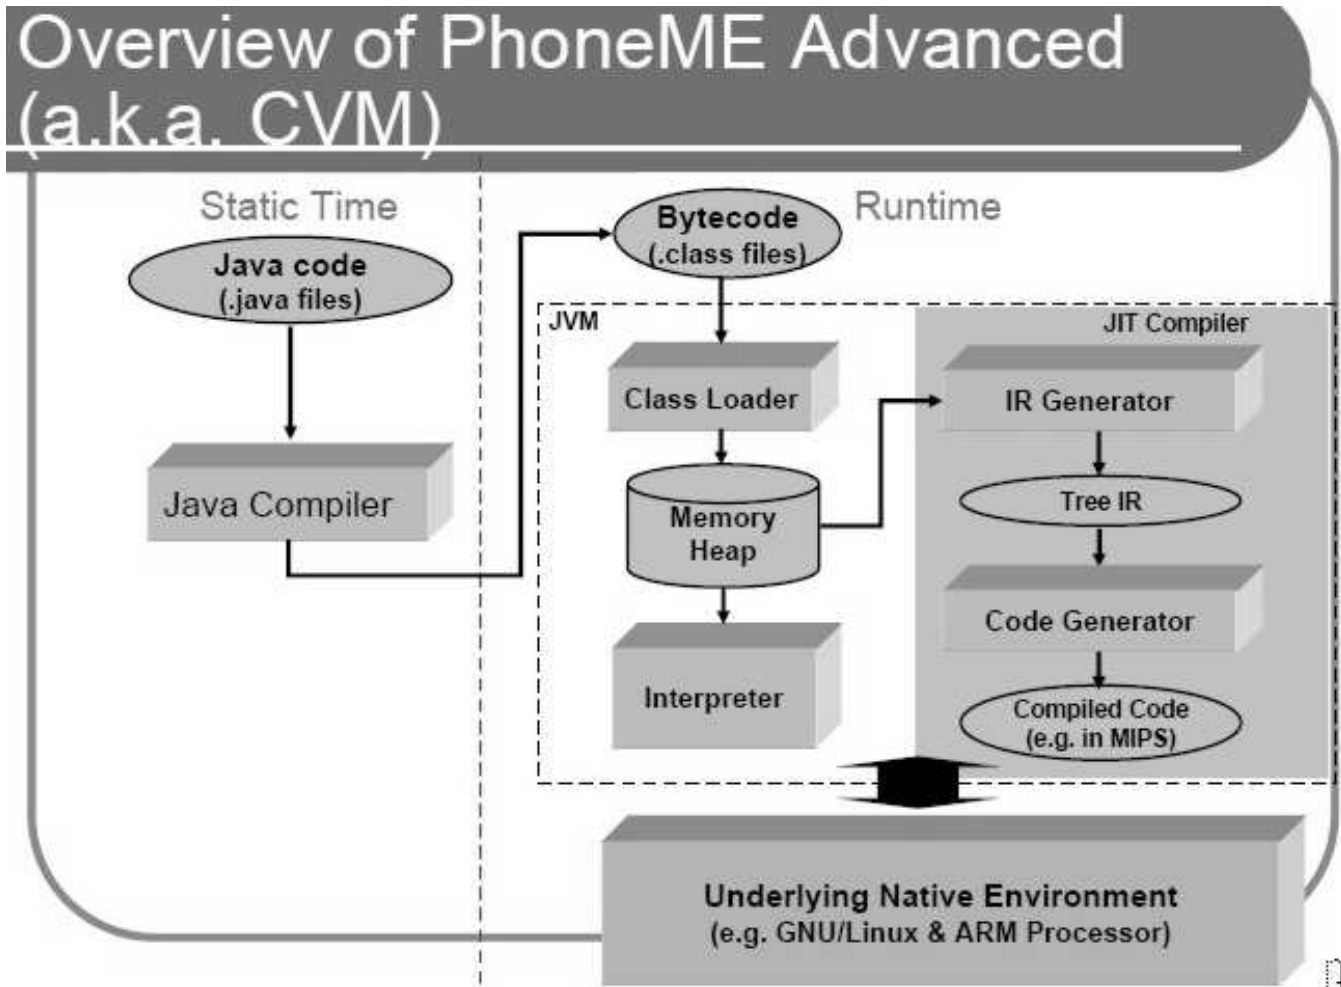


Figure 1: Java architecture

# 1 Java Virtual Machines

JVM is an abstract specification of a virtual machine intended for Java programs. See The Java Virtual Machine Specification, 2nd ed. by Tim Lindholm and Frank Yellin. There are many implementations of JVM on various hardware and software platforms, some of which are free.

When a Java application starts, a run-time instance of JVM is created. The instance is destroyed when the application completes. Each Java application runs inside its own JVM instance.

JVM is a stack-based machine. All arithmetic operations are done through the *operand stack*. Byte code is essentially a postfix form of the program.

JVM also does many security checks before running a procedure.

JVM makes use of dynamic loading and linking. It loads classes

during run time only when necessary.

JVM provides a *native method interface* for invoking native method libraries.

Run-time data structures of a JVM instance:

- a method area: shared by all threads in the JVM
- a heap: shared by all threads in the JVM
- a Java stack per thread
- a program counter per thread
- native method stacks




Figure 2: JVM run-time data structure

## *Data types in JVM*

- primitive type
  1. numeric type: float, double, byte, short, int, long, char
  2. boolean type
  3. return address type
- reference type
  1. class type
  2. interface type
  3. array type



## *class loader*

The class loader locates a class when reference, possibly from a remote machine, loads that class, and initializes that class.

- loading: find and import the class
- linking
  1. verification: ensure the imported class file is correct, not tampered.
  2. preparation: allocate memory for class variables and set up their default values.
  3. resolution: translate symbolic reference to direct references.
- initialization: initialize class variables to their starting values.

Every JVM includes a *bootstrap class loader*. The bootstrap class loader in Sun Java 2 SDK looks for system classes only in the installation directory. It will not consult the `classpath` environment variable.

On the other hand, the `system class loader` will consult `classpath`. The system class loader is a user-defined class loader that is created when the JVM starts up.

### *Method area*

Methods from the loaded classes are stored in the method area.

Class (static) variables are also stored in the method area.

All threads share the same method area.

### *Constant pool*

The constant pool is the symbol table for a class file. It contains literals (strings, integers, floating-point constants) and symbolic references to types, methods, and fields in the class.

Constant pool is used during dynamic linking.

All constant pools are allocated in the method area.

## *Heap*

All instances of a class or an array are allocated from the heap, which is shared among all the threads for the same application.

There is a separate heap for each individual running application.

JVM provides an instruction that allocates memory for new objects on the heap but does not provide instructions to free the memory. Instead JVM provides a garbage collector for reclaiming useless memory.

## *Garbage collection*

Pointers exist in the Java stacks, the heap, the method area, and the native method stacks.

## *Arrays*

In Java, arrays are objects with a class. They are stored in the heap.

Multi-dimensional arrays are represented as arrays of arrays. (See L6-jvm.pdf, p. 15.)

The program counter occupies one word. It can be a native pointer or an offset from the beginning of a method's bytecode.

## *Java stacks*

A Java stack consists of many frames. A frame includes local variables, operand stack, and frame data.

The size of a frame (in particular, the size of the operand stack) for a method is calculated by the compiler. This creates a constraint on methods.

The local variables in a frame are addressed by their indices. A method's parameters are stored in the local variable section in the frame. In the frame for an instance method, the first word is a hidden pointer to the object itself. There is not such a pointer in the frame for a class method.

JVM is essentially a stack machine. Many instructions make use of the operand stack, which is allocated in the stack frame for a method.

The frame data section includes data for constant pool resolution,

normal method returns, and exception handling.

Many instructions reference entries in the constant pool, such as pushing a constant value from the constant pool to the operand stack, referencing classes and arrays when creating new instances, accessing fields of a class, or invoking methods in the class, and determining the class or interface of a particular objects, etc.



Native methods also need a stack. This native stack is separated from the Java stack, which is intended for the Java methods.

Note that Java methods may invoke native methods. Similarly, native methods may also invoke Java methods as well as other native methods.

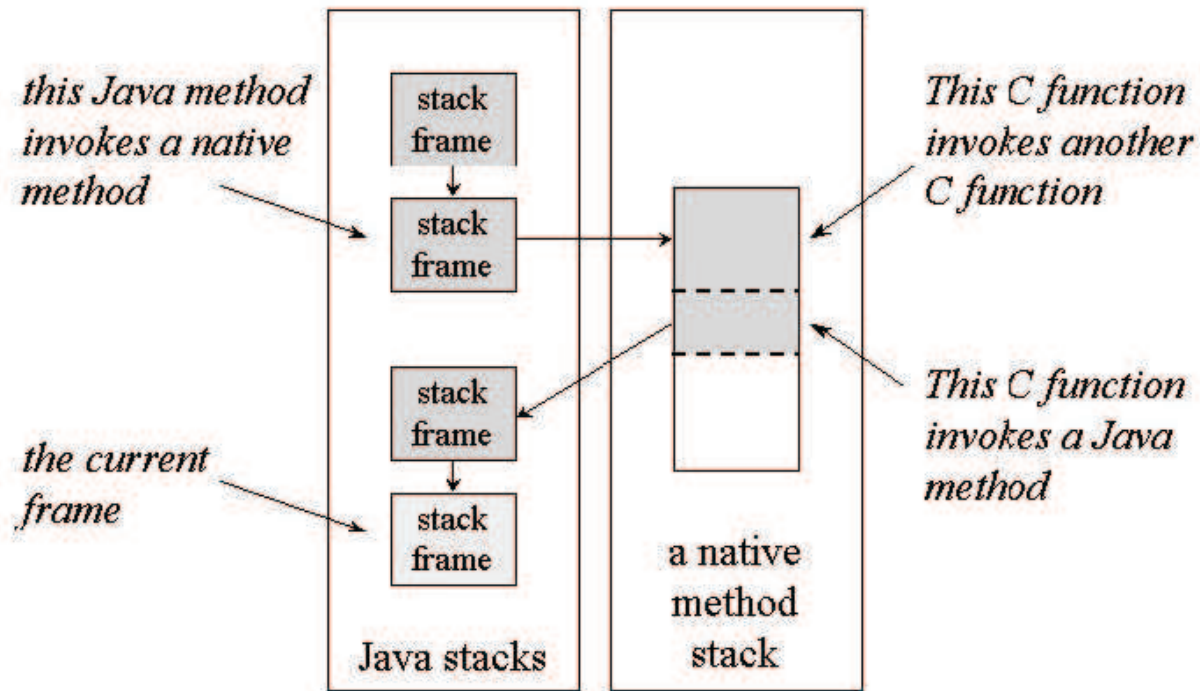


Figure 3: Java stacks vs. native stack.

A JVM may use invisible threads for other tasks, such as garbage collection.

Java instruction format:

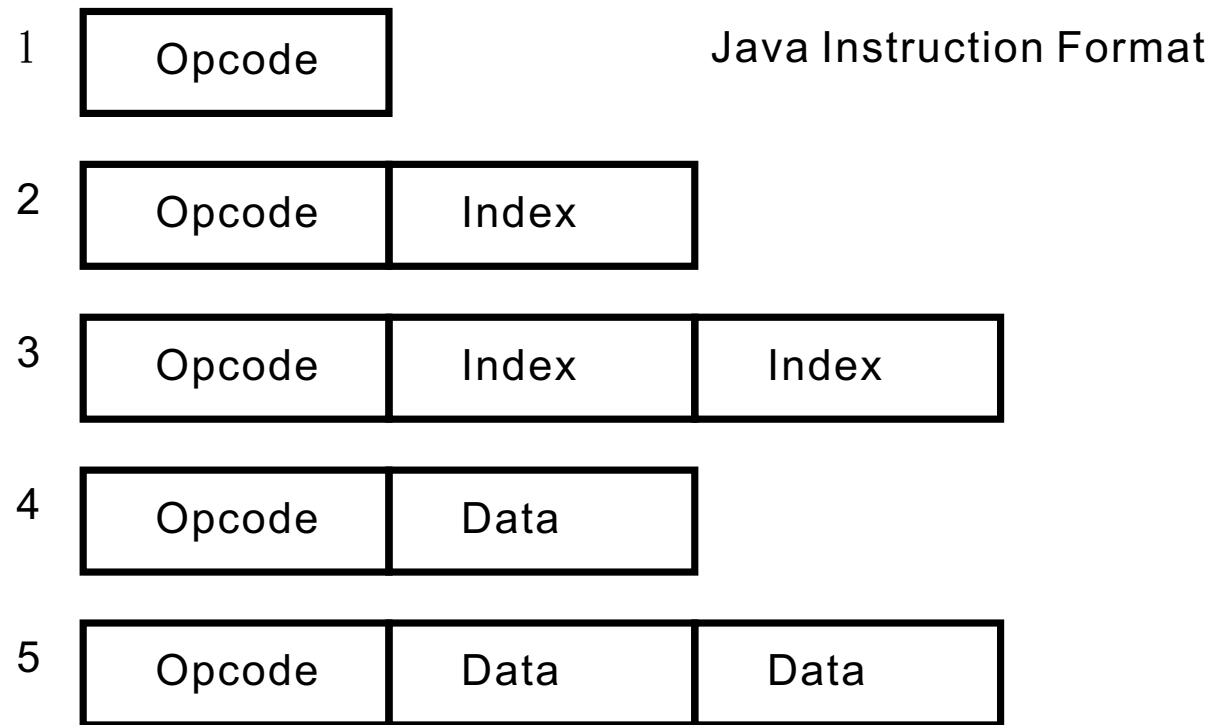


Figure 4: 6 Java instruction formats.

Java instruction set: see the manual.

- load and store instructions
- arithmetic instructions
- type conversion instructions
- object creation and manipulation instructions
- operand stack management instructions
- control transfer instructions
- method invocation instructions
- method return instructions



## 2 Bytecode

<b>Headers:</b> Magic number, version
<b>Constant Pool:</b> all external references; numerical and string constants; type signatures
<b>Class Data:</b> access flags, name, superclass, interfaces
<b>Fields:</b> indexes into constant pool for name and type; access flags; (optional) constant value
<b>Methods:</b> indexes into constant pool for name and type; access flags; code; exception handlers and thrown exceptions, debug information
<b>Attributes:</b> source file

Figure 1. Structure of a Java classfile

(slide 3) Operands for a bytecode instruction may be taken from

- the operand stack
- immediate operand (part of the instruction)
- implicit operand (in the instruction, e.g., `iload1`)
- local variable (from the run-time stack)
- the constant pool

(slide 18) Need to do bytecode verification to check the depth of the operand stack. Different paths from `<start>` to `<stop>` use the same amount of the operand stack.

## Constants, Local Variables, and Control Structures

<pre>void spin() {     int i;     for (i = 0; i &lt; 100; i++) {         ; // Loop body is empty     } }</pre>	<pre>Method void spin() 0  iconst_0  // Push constant 0 1  istore_1  // local var 1 is i 2  goto 8    // loop starts here 5  iinc 1 1   // i++ 8  iload_1    // Push local var 1 9  bipush 100 // Push constant 100 11 if_icmplt 5 // Compare and loop 14 return    // Return void</pre>
--	--

A new frame is created for every method invocation. Local variables are located in the current frame.

A frame also has an operand stack.

We use `iconst_0` (with an implicit operand, 1 byte) instead of



`bipush 0` (2 bytes).

The local variable `i` is stored in the current frame. We also use `iload_1` and `istore_1` to load from local variable 1 and store the top of the stack to local variable 1, respectively.

Note `iinc` increments a local variable, rather than the value on top of the operand stack. This saves a `push` and a `pop` operations. `iinc` is intended for loops.

For `int` type, we may use `if_icmplt` for comparison and jumping. For `double` type, we have to use two instructions: `dcmplt` and `iflt`. Java bytecode provides more support for `int` type than other types.

## Spins on Doubles

```
void dspin() {  
    double i;  
    for (i = 0.0; i < 100.0; i++) {  
        ; // Loop body is empty  
    }  
}
```

```
Method void dspin()  
0  dconst_0 // Push constant 0.0  
1  dstore_1  
   // store in local var 1 and 2  
2  goto 9  
5  dload_1 //Push var 1 and 2  
6  dconst_1 // Push constant 1.0  
7  dadd     // double-add  
8  dstore_1 // in var 1 and 2  
9  dload_1 //Push var 1 and 2  
10 ldc2_w #4 // Push 100.0  
13 dcmpg   // no if_dcmplt inst  
14 iflt 5 // Compare and loop  
17 return // Return void
```

A double value occupies two local variables.

```
double doubleLocals(double d1, double d2) {  
    Method double doubleLocals(double,double)  
    return d1 + d2;    0  dload_1    // 1st arg in local var 1 and 2  
}                      1  dload_3    // 2nd arg in local var 3 and 4  
                      2  dadd  
                      3  dreturn  
                      // The return value is on stack top.
```

There are no byte, short, and char versions of load, store or add instructions.

<pre>void sspin() {     short j;     for (j = 0; j &lt; 100; j ++) {         ;    // empty loop     } }</pre>	<pre>Method void sspin() 0  iconst_0 1  istore_1 // j = 0 2  goto 10 5  iload_1  // Treat short as int 6  iconst_1 7  iadd 8  i2s // Truncate int to short 9  istore_1 10 iload_1  // check if j &lt; 100 11 bipush 100 13 if_icmplt 5 16 return</pre>
---	--

`byte`, `short`, and `char` are transformed to the `int` type inside JVM.  
`byte` and `short` are sign-extended to `int`; `char` is zero-extended.  
Operations on `byte`, `short`, and `char` are done by `int` instructions.  
`long` and floating-point types are supported by JVM, lacking only conditional control-transfer instructions.

JVM generally does arithmetic operations on the operand stack (`iinc` is an exception).

```
int align2grain(int j, int grain) {  
    return ((j + grain-1) & ~(grain-1));  
}
```

Method `int align2grain(int,int)`

0	<code>iload_1</code>	// push j
1	<code>iload_2</code>	// push grain
2	<code>iadd</code>	// j + grain
3	<code>iconst_1</code>	// push int constant 1
4	<code>isub</code>	// j + grain - 1
5	<code>iload_2</code>	// Push grain
6	<code>iconst_1</code>	// Push int constant 1
7	<code>isub</code>	// Subtract; push result
8	<code>iconst_m1</code>	// Push int constant -1
9	<code>ixor</code>	// Do XOR; push result

```
10    iand
11    ireturn    // return value is on stack top.
```

## Run-Time Constant Pool

(constant pool in obfuscation)

Many numeric constants as well as objects, fields, and methods are accessed via the run-time constant pool of the current class.

Data of types `int`, `float` as well as references to `String` instances are managed with `ldc` and `ldc_w` instructions. `ldc_w` has a larger index range.

`ldc2_w` is for long and double.

`byte`, `short`, `int`, and `char` are managed with `bipush`, `sipush` and `iconst_<i>`. Certain small floating-point constants can be done with `fconst_<f>` and `dconst_<d>`.

<code>void useManyNumeric() {</code>	<code>Method void useManyNumeric()</code>
<code>    int k = 100;</code>	<code>0 bipush 100 //Push small int with bipush</code>
<code>    int j = 1000000;</code>	<code>2 istore_1</code>
<code>    long l1 = 1;</code>	<code>3 ldc #1 // Push int constant 1000000;</code>
	<code>// a larger int value uses ldc</code>



```

long l2 = 0xffffffff; 5 istore_2
double d = 2.2;      6 lconst_1
                      // A tiny long value uses short, fast lconst_1
}                    7 lstore_3
                      8 ldc2_w #6
                      // Push long 0xffffffff (that is, an int -1); any
                      // long constant value can be pushed using ldc2_w
                      11 lstore 5
                      13 ldc2_w #8
                      // Push double constant 2.200000; uncommon
                      // double values are also pushed using ldc2_w
                      16 dstore 7

```

## A Simple while Loop

```
void whileInt() {  
    int j = 0;  
    while (j < 100) {  
        j ++;  
    }  
}
```

```
Method void whileInt()  
0      iconst_0  
1      istore_1  
2      goto 8  
5      iinc 1 1  
8      iload_1  
9      bipush 100  
11     if_icmplt 5  
14     return
```

Note that the test is at the bottom of the loop. This incurs an extra `jump` instruction if the loop body is not executed. However, if the loop is executed at least once, this code saves one `jump` per iteration.

```

void whileDouble() {
    double k = 0.0;
    while (k < 100.1) {
        k ++;
    }
}

```

```

Method void whileDouble()
0      dconst_0    // 0
1      dstore_1    // k
2      goto 9
5      dload_1
6      dconst_1
7      dadd
8      dstore_1    // k++
9      dload_1     // k
10     ldc2_w #4
      // Push double constant 100.1
13     dcmpg
14     iflt 5
      // To do the compare and branch
      // we have to use two instructions
17     return

```



Each floating-point type has two comparison instructions: `fcompl` and `fcmpg` for type `float` and `dcmpl` and `dcmpg` for type `double`. The variants differ only in their treatment of NaN. NaN is unordered so all floating-point comparisons fail if any operand is NaN.

<code>int lessThan100(double d) {</code>	<code>Method int lessThan100(double)</code>
<code>if (d &lt; 100.0) {</code>	0 <code>dload_1</code> <code>// d</code>
<code>return 1;</code>	1 <code>ldc2_w #4</code> <code>// double 100.0</code>
<code>} else {</code>	4 <code>dcmpg</code>
	<code>// Push 1 if d is NaN or d &gt; 100.0;</code>
	<code>// push 0 if d == 100.0; -1 o.w.</code>
<code>return -1;</code>	5 <code>ifge 10</code> <code>// Branch on 0 or 1</code>
<code>}</code>	8 <code>iconst_1</code> <code>// 1</code>
<code>}</code>	9 <code>ireturn</code>
	10 <code>iconst_m1</code> <code>// -1</code>
	11 <code>ireturn</code>

## Receive Arguments

If  $n$  arguments are passed to an instance method, they are received in the local variables numbered 1 through  $n$ . By convention, an instance method is passed a **reference** to its instance in local variable 0 (that is, **this** in Java).

<code>int addTwo(int i, int j) {</code>	<code>Method int addTwo(int,int)</code>
<code>    return i + j;</code>	<code>0    iload_1 // Push local var 1, i</code>
<code>}</code>	<code>1    iload_2 // Push local var 2, j</code>
	<code>2    iadd    // Add</code>
	<code>3    ireturn // Return int result</code>

However, for `class` methods, there is not such a `reference` parameter. The usual parameters are at local variables 0, 1, etc.

<code>static int add2Static(int i, int j)</code>	<code>Method int add2Static(int,int)</code>
<code>{</code>	
<code>  return i + j;</code>	0 <code>iload_0</code> <code>// i</code>
<code>}</code>	1 <code>iload_1</code> <code>// j</code>
	2 <code>iadd</code>
	3 <code>ireturn</code>

There are 4 types of invoke instructions:

1. `invokevirtual index1 index2`: for invoking instance methods. *index1 index2* together forms an index into the constant pool where there is a descriptor of the called method (including address, number and types of arguments, max size of the operand stack).
2. `invokeinterface`: for interface methods
3. `invokespecial`: for special functions, such as initialization methods. The storage of an object is allocated with the `new` instruction and one of the constructor is called with the `invokespecial` instruction.
4. `invokestatic`: for class methods

There are several `return` instructions:

1. `ireturn`: returns an integer result



2. `dreturn`: returns a double result
3. `areturn`: returns an address result
4. `return`: returns a void value

## Call an Instance Method with `invokevirtual`

The `ireturn` instruction takes the `int` value returned by `addTwo`, on the operand stack of the current frame, and pushes it onto the operand stack of the frame of the invoker (of `add12and13`).

The argument to `invokevirtual`, i.e., `#4`, is a symbolic reference.

<code>int add12and13() {</code>	<code>Method int add12and13()</code>
<code>    return addTwo(12, 13);</code>	<code>0  aload_0  // Push local var 0, this</code>
<code>}</code>	<code>1  bipush 12  // Push int constant 12</code>
	<code>3  bipush 13  // Push int constant 13</code>
	<code>5  invokevirtual #4</code>
	<code>// Method Example.addTwo(II)I</code>
	<code>// signature</code>
	<code>// Return int on top of operand stack;</code>
	<code>// it is the int result of addTwo()</code>
	<code>8  ireturn</code>

## Call a Class Method with `invokestatic`

Note that there is no `this` reference in local variable 0.

```
int add12and13() {  
    return add2Static(12, 13);  
}
```

```
Method int add12and13()  
0  bipush 12  
2  bipush 13  
4  invokestatic #3  
  // Example.add2Static(II)I  
7  ireturn
```

## Call Instance Initialization Methods with `invokespecial`

`invokespecial` is also used to invoke methods in the superclass (`super`) and `private` methods. Note that methods called using `invokespecial` always pass `this` to the invoked method as its first argument. As usual, it is received in local variable 0.

```
class Near {  
    int it;  
    public int getItNear() {  
        return getIt();  
    }  
  
    private int getIt() {  
        return it;  
    }  
}  
  
Method int getItNear()  
0      aload_0  
1      invokespecial #5  
      // Method Near.getIt()I  
4      ireturn
```

```
class Far extends Near {  
    int getItFar() {  
        return super.getItNear();  
    }  
}
```

```
Method int getItFar()  
0      aload_0  
1      invokespecial #4  
      // Method Near.getItNear()I  
4      ireturn
```

## Create Instances with new

Once the class instance is created and its instance variables (including those of the class and all of its superclasses) have been initialized to their default values, an instance initialization method of the new class instance is invoked. Note that all constructors of a class are named `init` inside JVM.

<code>Object create() {</code>	<code>Method java.lang.Object create()</code>
<code>    return new Object();</code>	<code>0  new #1  // Class java.lang.Object</code>
<code>}</code>	<code>3  dup</code>
	<code>4  invokespecial #4</code>
	<code>// Method java.lang.Object.&lt;init&gt;()V</code>
	<code>7  areturn</code>

Class instances are passed and returned (as **reference** types) very much like numeric values, though with its own complement of instructions.

```
int i; // An instance var
MyObj example() {
    MyObj o = new MyObj();
    return silly(o);
}
```

```
Method MyObj example()
0      new #2    // Class MyObj
3      dup
4      invokespecial #5
// Method MyObj.<init>()V
7      astore_1  // o
8      aload_0   // this
9      aload_1   // o
10     invokevirtual #4
// Method Example.silly(LMyObj;)LMyObj;
13     areturn
```

```
MyObj silly(MyObj o) {  
    if (o != null) {  
        return o;  
    } else {  
        return o;  
    }  
}
```

```
Method MyObj silly(MyObj)  
0    aload_1    // o  
1    ifnull 6  
4    aload_1  
5    areturn  
6    aload_1  
7    areturn
```



The fields of an instance (i.e., instance variables) are accessed with `getfield` and `putfield`.

```
int i;  // instance var
```

```
...
```

```
void setIt(int value) {  
    i = value;  
}
```

```
Method void setIt(int)
```

```
0    aload_0      // this
```

```
1    iload_1      // value
```

```
2    putfield #4  // Field Example.i
```

```
5    return       // void
```

```
int getIt() {  
    return i;  
}
```

```
Method int getIt()
```

```
0    aload_0      // this
```

```
1    getfield #4  // Field Example.i
```

```
4    ireturn      // int
```

Some data movement instructions

`iconst_1`: pushes the integer constant 1 onto the stack.

`bipush data`: pushes a one-byte signed integer

`sipush data1 data2`: pushes a two-byte signed integer

`ldc index`: push single-word constant onto stack

`ldc_w index1 index2`: push single-word constant onto stack (wide index)

**Conditional branches** make use of PC relative addresses.

`ifeq data1 data1`: check if the integer on stack top is zero. If so, branch to the PC relative address with an offset calculated by concatenating the two data bytes.

`if_icmpeq data1 data2`: check if the two integers on stack top are equal. If so, branch to the PC relative address with an offset calculated by concatenating the two data bytes.

`ifnull data1 data2`: check the reference on stack top is null.

`newarray` is used to create arrays of numeric types.

```
void createBuffer() {  
    int buffer[];  
    int bufsz = 100;  
    int value = 12;  
    buffer = new int[bufsz];  
    buffer[10] = value;  
    value = buffer[11];  
}
```

```
Method void createBuffer()  
0      bipush 100  
2      istore_2    // bufsz  
3      bipush 12  
5      istore_3    // value  
6      iload_2     // Push bufsz  
7      newarray int // create array  
9      astore_1    // buffer  
10     aload_1     // Push buffer  
11     bipush 10  
13     iload_3     // Push value  
14     iastore     // buffer[10]  
15     aload_1     // Push buffer  
16     bipush 11  
18     iaload      // buffer[11]
```

```
19     istore_3    // value
20     return
```

`anewarray` is used to create 1-dim arrays of object references.

<code>void createThreadArray() {</code>	<code>Method void createThreadArray()</code>
<code>    Thread threads[];</code>	<code>0    bipush 10    // Push 10</code>
<code>    int count = 10;</code>	<code>2    istore_2    // count</code>
<code>    threads = new Thread[count];</code>	<code>3    iload_2</code>
<code>    threads[0] = new Thread();</code>	<code>4    anewarray class #1</code>
	<code>// Create array of class Thread</code>
<code>}</code>	<code>7    astore_1    // threads</code>
	<code>8    aload_1</code>
	<code>9    iconst_0</code>
	<code>10   new #1</code>
	<code>// Create instance of class Thread</code>
	<code>13   dup</code>
	<code>14   invokespecial #5</code>
	<code>// pass to instance initialization method</code>
	<code>// Method java.lang.Thread.&lt;init&gt;()V</code>

```
17     astore
// Store new Thread at threads[0]
18     return
```

`multianewarray` creates multi-dimensional arrays.

<code>int[][][] create3DArray() {</code>	<code>Method int create3DArray() [] [] []</code>
<code>    int grid[][][];</code>	<code>0    bipush 10</code>
<code>    grid = new int[10][5][];</code>	<code>2    iconst_5</code>
<code>    return grid;</code>	<code>3    multianewarray #1 dim #2</code>
<code>        // Class [[[I, a three dimensional int array;</code>	
<code>        // only create first two dimensions</code>	
<code>}</code>	<code>7    astore_1 // grid</code>
	<code>8    aload_1</code>
	<code>9    areturn</code>

All arrays have associated lengths, by the `arraylength` instruction.



tableswitch and lookupswitch instructions:

int chooseNear(int i) {	Method int chooseNear(int)
switch (i) {	0    iload_1    // local var i
case 0: return 0;	1    tableswitch 0 to 2:
	0: 28
	1: 30
	2: 32
	default: 34
case 1: return 1;	28    iconst_0
case 2: return 2;	29    ireturn
default: return -1;	30    iconst_1
}	31    ireturn
}	32    iconst_2
	33    ireturn
	34    iconst_m1
	35    ireturn



`tableswitch` and `lookupswitch` instructions operate on the `int` type.

When the cases of `switch` are sparse, the table representation of `tableswitch` becomes inefficient in terms of space. `lookupswitch` may be used instead.

<code>int chooseFar(int i) {</code>	<code>Method int chooseFar(int)</code>
<code>    switch (i) {</code>	<code>0    iload_1 // local var i</code>
<code>        case -100: return -1;</code>	<code>1    lookupswitch 3:</code>
	<code>        -100: 36</code>
	<code>        0: 38</code>
	<code>        100: 40</code>
	<code>        default: 42</code>
<code>        case 0:    return 0;</code>	<code>36    iconst_m1</code>
<code>        case 100: return 1;</code>	<code>37    ireturn</code>
<code>        default:  return -1;</code>	<code>38    iconst_0</code>
<code>    }</code>	<code>39    ireturn</code>

}

40	iconst_1
41	ireturn
42	iconst_m1
43	ireturn

## Operations on Operand Stack

<pre>public long nextIndex() {     return index++; } private long index = 0;     // One copy of this is consumed     // index above the original this</pre>	<pre>Method long nextIndex() 0    aload_0  // Push this 1    dup      // two copies 2    getfield #4    // index 5    dup2_x1 6    lconst_1  // long constant 1 7    ladd 8    putfield #4  // index 11   lreturn     // use the original copy of this</pre>
---	--

## Throw Exceptions With `throw`

```
void cantBeZero(int i) throws TestExc {  
    if (i == 0) {  
        throw new TestExc();  
    }  
}  
  
Method void cantBeZero(int)  
0    iload_1    // i  
1    ifne 12    // if i != 0  
4    new #1     // instance of TestExc  
7    dup  
    // One reference is for constructor  
8    invokespecial #7  
    // Method TestExc.<init>()V  
11   athrow  
    // Second reference is thrown  
12   return
```

## Compiling try-catch Blocks

```
void catchOne() {  
    try {  
        tryItOut();  
    } catch (TestExc e) {  
        handleExc(e);  
    }  
}
```

```
Method void catchOne()  
0  aload_0  // this  
1  invokevirtual #6  
   // Method Example.tryItOut()V  
4  return   // normal return  
5  astore_1 // 1st handler  
   // Store thrown value in local var 1  
6  aload_0  // Push this  
7  aload_1  // Push thrown value  
8  invokevirtual #5 // handler method:  
   // Example.handleExc(LTestExc;)V  
11 return   // Return after TestExc
```

Exception table:

From	To	Target	Type
0	4	5	Class TestExc

The `catch` block generates an entry in the exception table. The exception shows that exceptions occurring between line 0 and line 3 (not including line 4) will be handled by the handler.

There could be multiple `catch` clauses.

There could also be nested `try-catch` blocks.



Here is an example with two exception handlers.

```
void catchTwo() {  
    try {  
        tryItOut();  
    } catch (TestExc1 e) {  
        handleExc(e);  
    } catch (TestExc2 e) {  
        handleExc(e);  
    }  
}
```

```
Method void catchTwo()  
0  aload_0  // this  
1  invokevirtual #5  
   // Method Example.tryItOut()V  
4  return   // normal return  
5  astore_1 // 1st handler  
6  aload_0  // push this  
7  aload_1  // push e  
8  invokevirtual #7  
   // Example.handleExc(LTestExc1;)V  
11 return // after handling TestExc1  
12 astore_1 // 2nd handler  
13 aload_0  // Push this  
14 aload_1  // Push e  
15 invokevirtual #7
```

```
18  return // after handling TestExc2
```

Exception table:

From	To	Target	Type
0	4	5	Class TestExc1
0	4	12	Class TestExc2

## Nested try-catch Blocks

```
void nestedCatch() {  
    try {  
        try {  
            tryItOut();  
        } catch (TestExc1 e) {  
            handleExc1(e);  
        }  
    } catch (TestExc2 e) {  
        handleExc2(e);  
    }  
}
```

```
Method void nestedCatch()  
0  aload_0  // push this  
1  invokevirtual #8  
   // Method Example.tryItOut()V  
4  return  // normal return  
5  astore_1 // 1st handler  
6  aload_0  // Push this  
7  aload_1  // Push e  
8  invokevirtual #7  
   // Example.handleExc1(LTestExc1;)V  
11 return//after handling TestExc1  
12 astore_1 // 2nd handler  
13 aload_0  // Push this  
14 aload_1  // Push thrown value  
15 invokevirtual #6
```

```
// Example.handleExc2(LTestExc2;)V
18  return//after handling TestExc2
```

Exception table:

From	To	Target	Type
0	4	5	Class TestExc1
0	12	12	Class TestExc2

Note the order of the two exception handlers in the exception table.

## Compiling finally

```
void tryFinally() {  
    try {  
        tryItOut();  
    } finally {  
        wrapItUp();  
    }  
}
```

```
Method void tryFinally()  
0  aload_0 // push this  
1  invokevirtual #6  
   // Method Example.tryItOut()V  
4  jsr 14  // call finally block  
7  return  
8  astore_1 // handler for any throw  
9  jsr 14  // Call finally block  
12 aload_1 // Push thrown value  
13 athrow // rethrow the value to invoker  
14 astore_2 // Beginning of finally,  
           // return address  
15 aload_0 // Push this  
16 invokevirtual #5  
   // Method Example.wrapItUp()V
```

```
19  ret 2 // Return from finally block
```

Exception table:

From	To	Target	Type
0	4	8	any

The `finally` block is compiled as an embedded subroutine and invoked with `jsr`. When `finally` block completes, `ret 2` returns control to the instruction following `jsr`. Specifically, `jsr` pushes the return address onto operand stack. `astore_2` saves that address in local variable 2. Finally, `ret 2` jumps to the address saved in local variable 2.

```

void tryCatchFinally() {
    try {
        tryItOut();
    } catch (TestExc e) {
        handleExc(e);
    } finally {
        wrapItUp();
    }
}

```

```

Method void tryCatchFinally()
0  aload_0
1  invokevirtual #4 // tryItOut()V
4  goto 16
7  astore_3 // Beginning of 1st handler
8  aload_0 // Push this
9  aload_3 // Push thrown value
10 invokevirtual #6 // handleExc
13 goto 16
16 jsr 26 // Call finally block
19 return // after handling TestExc
20 astore_1 // Beginning of 2nd handler
21 jsr 26 // Call finally block
24 aload_1 // Push thrown value...
25 athrow //...and rethrow to invoker
26 astore_2 // Beginning of finally

```

```
27  aload_0  // Push this
28  invokevirtual #5 // wrapItUp()V
31  ret 2     // Return from finally
```

Exception table:

From	To	Target	Type
0	4	7	Class TestExc
0	16	20	any



## Synchronization with `monitorenter` and `monitorexit`

```
void onlyMe(Foo f) {  
    synchronized(f) {  
        doSomething();  
    }  
}
```

Method `void onlyMe(Foo)`

```
0  aload_1  // Push f  
1  astore_2 // Store it in local var 2  
2  aload_2  // Push local variable 2 (f)  
3  monitorenter // Enter the monitor of f  
4  aload_0  // this  
5  invokevirtual #5 // doSomething()V  
8  aload_2  // Push local variable 2 (f)  
9  monitorexit // Exit the monitor of f  
10 return    // Return normally  
11 aload_2  // In case of throw, stop here  
12 monitorexit // Be sure to exit monitor  
13 athrow    // ...then rethrow to invoker
```

Exception table:

From	To	Target	Type
------	----	--------	------

4      8      11      any

However, the most common use of synchronization is a **synchronized** method. A **synchronized** method is *not* implemented with **monitorenter** and **monitorexit**. Rather, an **ACC\_SYNCHRONIZED** flag of the method is set, which is checked by the method invocation instructions. When invoking a method with the **ACC\_SYNCHRONIZED** flag set, the current thread acquires a monitor, invokes the method, and releases the monitor after the method completes. If an exception occurs during a method execution and the method does not handle the exception, the monitor of the method is automatically released before the exception is thrown out of the **synchronized** method.

## **Nested (i.e. Inner) Classes and Interfaces**

See inner classes specification.

### 3 Sample Java Code

```
public static void main(String[] args){  
    int i;  
    for(i = 0; i < 10; i++) {  
        if (i == 5) System.out.println("i = 5");  
        else System.out.println("i != 5");  
    }  
}
```

```
0:   iconst_0  
1:   istore_1  
2:   iload_1  
3:   bipush  10  
5:   if_icmpge 38  
8:   iload_1
```

```
9:   iconst_5
10:   if_icmpne 24
13:   getstatic #2;
       //Field java/lang/System.out:Ljava/io/PrintStream;
16:   ldc      #3; //String i = 5
18:   invokevirtual #4;
       //Method java/io/PrintStream.println:(Ljava/lang/String;)V
21:   goto     32
24:   getstatic #2;
       //Field java/lang/System.out:Ljava/io/PrintStream;
27:   ldc      #5;   //String i != 5
29:   invokevirtual #4;
       //Method java/io/PrintStream.println:(Ljava/lang/String;)V
32:   iinc     1, 1
35:   goto     2
38:   return
```

